

Homework 1

Joshua Michalenko
ELEC 677: Deep Learning
Dr. Ankit Patel

10/04/16

1 Problem 1: Backpropagation in a simple Nural Network

1.1 Part A: Dataset

The plot of the Two Moons dataset is displayed in Figure 1

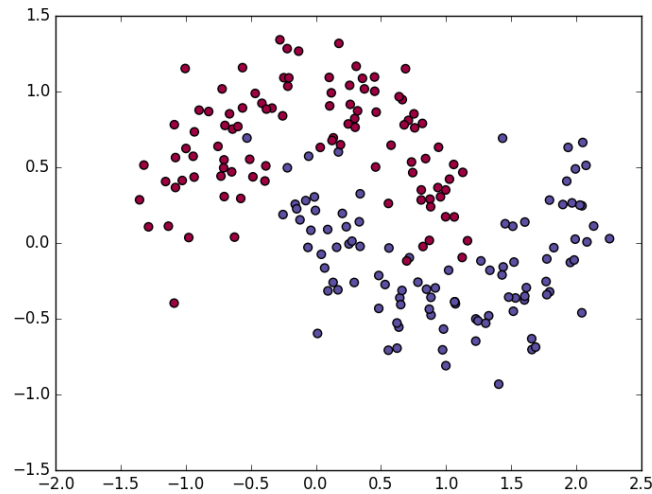


Figure 1: Two moons dataset displayed in matplotlib

1.2 Part B - Activation Functions

1.2.1 Part B1 - Implement activation functions

If you look in my code you can see I implemented the activations function.

1.2.2 Part B2 - Derivative Derivations

Part B2.1 - Derivative for Sigmoid function

$$\begin{aligned}\sigma(z) = f(z) &= \frac{1}{1 + e^{-z}} = (1 + e^{-z})^{-1} \\ \frac{df}{dz} &:= -1 * (1 + e^{-z})^{-2} * -e^{-z} = e^{-z}(1 + e^{-z})^{-2} \quad (\text{by chain rule}) \\ &= \frac{e^{-z}}{(1 + e^{-z})^2} \\ &= \frac{1 - e^{-z} + 1}{(1 + e^{-z})^2} \quad (\text{add and subtract a 1}) \\ &= \frac{1 + e^{-z}}{(1 + e^{-z})^2} - \frac{1}{(1 + e^{-z})^2} \quad (\text{split terms}) \\ &= \frac{1}{1 + e^{-z}} - \frac{1}{(1 + e^{-z})^2} \\ &= \sigma(z) - \frac{1}{(1 + e^{-z})^2} \\ &= \sigma(z) - \left(\frac{1}{(1 + e^{-z})} \right)^2 \\ &= \sigma(z) - \sigma(z)^2 \\ &\boxed{f'(z) = \sigma(z)(1 - \sigma(z))}\end{aligned}$$

Part B2.2 - Derivative for Tanh function

$$\begin{aligned}f(z) = \tanh(z) &= \frac{\sinh(z)}{\cosh(z)} \\ &= \frac{e^z - e^{-z}}{e^z + e^{-z}} \\ \frac{df}{dz} &:= \frac{\cosh(z) * \sinh'(z) - \sinh(z) * \cosh'(z)}{(\cosh(z))^2} \quad (\text{by quotient rule}) \\ &= \frac{\cosh(z)^2 - \sinh(z)^2}{(\cosh(z))^2} \\ &= \frac{\cosh(z)^2}{\cosh(z)^2} - \frac{\sinh(z)^2}{\cosh(z)^2} \\ &\boxed{f'(z) = 1 - \tanh^2(z)}\end{aligned}$$

Part B2.3 - Derivative for ReLU function

$$\begin{aligned}\text{ReLU}(z) = f(z) &= \max(0, z) \\ \frac{df}{dz} &:= \begin{cases} 0 & z \leq 0 \\ 1 & z \geq 0 \end{cases} \quad (\text{I think the derivative here is pretty obvious})\end{aligned}$$

1.2.3 Part B3 - Implement activation function gradients

If you look in my code you can see I implemented the activations function gradients.

1.3 Part C - Build the Neural Network

If you look in my code you can see I implemented the feedforward and loss functions

1.4 Part D - Backpropagation Derivations

1.4.1 Derivations of Backpropagation Equations

Keep in mind that the process for a single feedforward pass of the network is defined by the following equations with $x \in \mathbb{R}^p$ being a single data sample with p features, $\mathbf{W}_1 \in \mathbb{R}^{n_1 \times p}$ is a weight matrix transitioning from the input later with p features to the number of units in hidden layer 1 n_1 , $b_1 \in \mathbb{R}^{n_1}$ is a bias vector, $z_1 \in \mathbb{R}^{n_1}$ are a vector of potentials for layer 1, $a_1 \in \mathbb{R}^{n_1}$ the corresponding activations, and \hat{y} are the resulting probabilities. $\mathbf{W}_2 \in \mathbb{R}^{n_2 \times n_1}$, $a_1 \in \mathbb{R}^{n_1}$, $z_2 \in \mathbb{R}^{n_2}$

$$\begin{aligned} z_1 &= \mathbf{W}_1 x + b_1 \\ a_1 &= \text{actFun}(z_1) \\ z_2 &= \mathbf{W}_2 a_1 + b_2 \\ a_2 &= \hat{y} = \text{softmax}(z_2) \end{aligned}$$

Where the softmax function is given by:

$$\text{softmax}(\mathbf{z})_c = \frac{e^{z_c}}{\sum_{d=1}^C e^{z_d}} = \frac{e^{z_c}}{\Sigma_C} \quad \text{for } c = 1 \dots C$$

With a cross entropy loss function of :

$$L(y, \hat{y}) = -\frac{1}{N} \sum_{n \in N} \sum_{j \in C} y_{n,j} \log(\hat{y}_{n,j})$$

With y being a one hot vector of the correct label, N being the number of training examples and C being the number of classes.

For the following derivations, the derivative of the softmax $\frac{\partial \hat{y}_i}{\partial z_j}$ is useful to know. It is calculated below. It is the derivative of \hat{y} with respect to one of the elements of the input vector z .

$$\begin{aligned}
\hat{y} = \text{softmax}(\mathbf{z})_c &= \frac{e^{z_c}}{\sum_{d=1}^C e^{z_d}} = \frac{e^{z_c}}{\Sigma_C} \quad \text{for } c = 1 \dots C \\
\text{if } i = j : \quad \frac{\partial y_i}{\partial z_i} &= \frac{\partial \frac{e^{z_i}}{\Sigma_C}}{\partial z_i} \\
&= \frac{e^{z_i} \Sigma_C - e^{z_i} e^{z_i}}{\Sigma_C^2} \\
&= \frac{e^{z_i}}{\Sigma_C} \frac{\Sigma_C - e^{z_i}}{\Sigma_C} \\
&= \frac{e^{z_i}}{\Sigma_C} \left(1 - \frac{e^{z_i}}{\Sigma_C}\right) \\
&= y_i(1 - y_i) \\
\text{if } i \neq j : \quad \frac{\partial y_i}{\partial z_j} &= \frac{\partial \frac{e^{z_i}}{\Sigma_C}}{\partial z_j} \\
&= \frac{0 - e^{z_i} e^{z_j}}{\Sigma_C^2} \\
&= -\frac{e^{z_i}}{\Sigma_C} \frac{e^{z_j}}{\Sigma_C} \\
&= -y_i y_j
\end{aligned}$$

The derivative of the cross entropy loss with respect to the second layer inputs are also a useful quantity to calculate, it is derived as follows. Note that I used the cross entropy for one example. Taking out the summation here makes things easier but I'll add it back in later.

$$\begin{aligned}
\frac{\partial L}{\partial \hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_{2i}} &= \frac{\partial L}{\partial z_{2i}} = -\sum_{j=1}^C \frac{\partial y_j \log(\hat{y}_j)}{\partial z_i} \\
&= -\sum_{j=1}^C y_j \frac{\partial \log(\hat{y}_j)}{\partial z_i} \\
&= -\sum_{j=1}^C y_j \frac{1}{\hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_i} \\
&= -\frac{y_i}{\hat{y}_i} \frac{\partial \hat{y}_i}{\partial z_i} - \sum_{j \neq i}^C \frac{y_j}{\hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_i} \quad \text{substitution from above derivation} \\
&= -\frac{y_i}{\hat{y}_i} \hat{y}_i (1 - \hat{y}_i) - \sum_{j \neq i}^C \frac{y_j}{\hat{y}_j} (-\hat{y}_j \hat{y}_i) \\
&= -y_i + y_i \hat{y}_i + \sum_{j \neq i}^C y_j \hat{y}_i \\
&= -y_i + \sum_{j=1}^C y_j \hat{y}_i \\
&= -y_i + \hat{y}_i \sum_{j=1}^C y_j \\
&= \hat{y}_i - y_i
\end{aligned}$$

$\frac{dL}{dW_2}$ **Derivation** We can rewrite the loss function with substituting in the feed-forward equations above as follows.

$$L(y, \hat{y}) = -\frac{1}{N} \sum_{n \in N} \sum_{j \in C} y_{n,j} \log(\hat{y}_{n,j})$$

$$\frac{\partial L}{\partial W_2} := \frac{\partial L}{\partial \hat{y}_{n,j}} \frac{\partial \hat{y}_{n,j}}{\partial z_2} \frac{\partial z_2}{\partial W_2}$$

Since j -th element of z_2 is given by:

$$z_{2_j} = \sum_{k=0}^{n_1} w_{j,k} a_{1_k} + b_{2_j}$$

The derivative of the j -th element of z_2 w.r.t. weight $w_{j,k}$ is then:

$$\frac{\partial z_{2_j}}{\partial w_{j,k}} = a_{1_k}$$

Using the derivation above for $\frac{\partial L}{\partial z_{2_j}}$

$$\frac{\partial L}{\partial w_{2_j,k}} = (\hat{y}_j - y_j) a_{1_k}$$

\therefore

$$\frac{dL}{dW_2} = (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{a}_1^T \in \mathbb{R}^{n_2 \times n_1}$$

The above equation only takes into consideration on data point
We extend this to multiple samples by plugging in the summation we originally took out.

$$\boxed{\frac{dL}{dW_2} = \frac{1}{N} \sum_{n \in N} (\hat{\mathbf{y}}_n - \mathbf{y}_n) \mathbf{a}_1^T \in \mathbb{R}^{n_2 \times n_1}}$$

$\frac{dL}{db_2}$ **Derivation** This derivation will be fairly similar to the previous section except now the gradient looks as follows:

$$\frac{\partial L}{\partial b_2} := \frac{\partial L}{\partial \hat{y}_{n,j}} \frac{\partial \hat{y}_{n,j}}{\partial z_2} \frac{\partial z_2}{\partial b_2}$$

Since j -th element of z_2 is given by:

$$z_{2j} = \sum_{k=0}^{n_1} w_{j,k} a_{1k} + b_{2j}$$

The derivative of the j -th element

of z_2 w.r.t. bias b_{2k} is then:

$$\frac{\partial z_{2j}}{\partial b_{2k}} = 1$$

Using the derivation above for $\frac{\partial L}{\partial z_{2j}}$

$$\frac{\partial L}{\partial b_{2k}} = (\hat{y}_j - y_j)$$

\therefore

$$\frac{dL}{db_2} = \hat{\mathbf{y}} - \mathbf{y} \in \mathbb{R}^{n_2}$$

The above equation only takes into consideration on data point.

We extend this to multiple samples by plugging in the summation we originally took out.

$$\boxed{\frac{dL}{db_2} = \frac{1}{N} \sum_{n \in N} \hat{\mathbf{y}}_{\mathbf{n}} - \mathbf{y}_{\mathbf{n}} \in \mathbb{R}^{n_2}}$$

$\frac{dL}{dW_1}$ **Derivation** We add one more layer to our previous derivation of $\frac{dL}{dW_2}$ for this derivation. The derivative can be expanded as follows

$$\frac{\partial L}{\partial W_1} := \frac{\partial L}{\partial \hat{y}_{n,j}} \frac{\partial \hat{y}_{n,j}}{\partial z_2} \frac{\partial z_2}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial W_1}$$

So we need to figure out the term $\frac{\partial z_2}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial W_1}$ by piecing to out. The j -th element of z_1 is given by:

$$z_{1j} = \sum_{i=0}^p w_{1j,i} x_i + b_{1j}$$

So the derivative w.r.t. $w_{1j,i}$ is then:

$$\frac{\partial z_{1j}}{\partial w_{1j,i}} = x_i$$

We know from part b that $\frac{\partial a_1}{\partial z_1}$ is dependent on the type of activation function but we derived all the derivatives of the three activation functions above so stating the derivatives is redundant. See the above part B.

$$\frac{\partial a_1}{\partial z_1} = \text{actFun}'(z_1) \quad (\text{See part B above})$$

The last part we need to derive is the $\frac{\partial z_2}{\partial a_1}$ term.

Since the j -th element of z_2 is given by:

$$z_{2_j} = \sum_{k=0}^{n_1} w_{2_j,k} a_{1_k} + b_{2_j}$$

we can conclude that the derivative w.r.t. a_{1_k} is given by

$$\frac{\partial z_{2_j}}{\partial a_{1_k}} = w_{2_j,k}$$

From above, we know that the other partial derivatives are

$$\frac{\partial L}{\partial \hat{y}_j} \frac{\partial \hat{y}_j}{\partial z_{2_i}} = \frac{\partial L}{\partial z_{2_i}} = \hat{y}_i - y_i$$

So if we piece this all together we come up with:

$$\frac{\partial L}{\partial w_{1_k,p}} := \frac{\partial L}{\partial \hat{y}_{n,c}} \frac{\partial \hat{y}_{n,c}}{\partial z_{2_c}} \frac{\partial z_{2_c}}{\partial a_{1_k}} \frac{\partial a_{1_k}}{\partial z_{1_k}} \frac{\partial z_{1_k}}{\partial w_{1_k,p}}$$

Which I changed some of the indexing to make the most sense. In this format $n \in N$, $c \in C$, $k \in \{0, \dots, n_1 - 1\}$, $p \in \{0, \dots, P - 1\}$ with P input features. The resulting derivative with respect to one weight in the first layer $w_{1_k,p}$ is then:

$$\frac{\partial L}{\partial w_{1_k,p}} = \sum_{c \in C} (\hat{y}_c - y_c) w_{2_c,k} \text{actFun}'(z_{1_k}) x_p$$

and the extention to matrix form is then

$$\frac{\partial L}{\partial \mathbf{W}_1} = (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{W}_2 \odot \text{actFun}'(\mathbf{z}_1) \mathbf{x}^T \in \mathbb{R}^{n_1 \times p}$$

If there is more than one sample (which there always is)

then we average over all of the samples and the final solution is then

$$\boxed{\frac{\partial L}{\partial \mathbf{W}_1} = \frac{1}{N} \sum_{n \in N} (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{W}_2 \odot \text{actFun}'(\mathbf{z}_1) \mathbf{x}^T \in \mathbb{R}^{n_1 \times p}}$$

$\frac{dL}{db_1}$ **Derivation** The derivation is very similar to the $\frac{dL}{dW_1}$ except the last part of the chain of partial derivatives is changed. The gradient equation of the Loss function with respect to a single bias term in the first layer is

$$\frac{\partial L}{\partial b_{1_k}} := \frac{\partial L}{\partial \hat{y}_{n,c}} \frac{\partial \hat{y}_{n,c}}{\partial z_{2_c}} \frac{\partial z_{2_c}}{\partial a_{1_k}} \frac{\partial a_{1_k}}{\partial z_{1_k}} \frac{\partial z_{1_k}}{\partial b_{1_k}}$$

So we only need to calculate $\frac{\partial z_{1_k}}{\partial b_{1_k}}$ first. Let's look at how the k th value is calculated.

$$z_{1_k} = \sum_{i=0}^p w_{1_k,i} x_i + b_{1_k}$$

So the derivative w.r.t. b_{1_k} is then:

$$\frac{\partial z_{1_k}}{\partial b_{1_k}} = 1$$

The partial derivative $\frac{\partial z_{1_k}}{\partial b_{1_k}} = 1$ simplified the equation for the partial derivative of the loss function w.r.t. b_{1_k} as

$$\frac{\partial L}{\partial b_{1_k}} = \sum_{c \in C} (\hat{y}_c - y_c) w_{2_{c,k}} \text{actFun}'(z_{1_k})$$

and the resulting matrix form is then

$$\frac{\partial L}{\partial \mathbf{b}_1} = (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{W}_2^T \text{actFun}'(\mathbf{z}_1) \in \mathbb{R}^{n_1}$$

If there is more than one sample (which there always is) then we have to sum up and average over all samples and the gradient becomes

$$\boxed{\frac{\partial L}{\partial \mathbf{b}_1} = \frac{1}{N} \sum_{n \in N} (\hat{\mathbf{y}} - \mathbf{y}) \mathbf{W}_2^T \odot \text{actFun}'(\mathbf{z}_1) \in \mathbb{R}^{n_1}}$$

1.4.2 Implementation of Derivations of Backpropogation Equations

If you look in my code you can see I implemented the backpropogation equations correctly.

1.5 Time to Have Fun - Training!

The decision boundaries are shown for tanh, sigmoid, and relu functions for 3, 10 and 20 hidden hidden units. Using the default 3 hidden units with different activation fuctions, we get Figures 2,3, and 4.

Increasing the number of hidden units to 10 gives us figure 5, 6 and 7

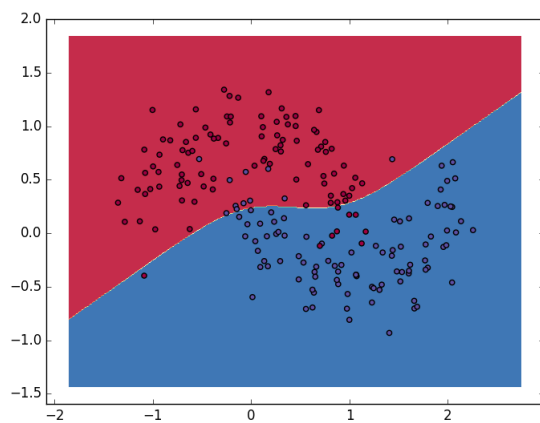


Figure 2: Default settings, 3 hidden units with tanh activation fuctions. Loss = .260

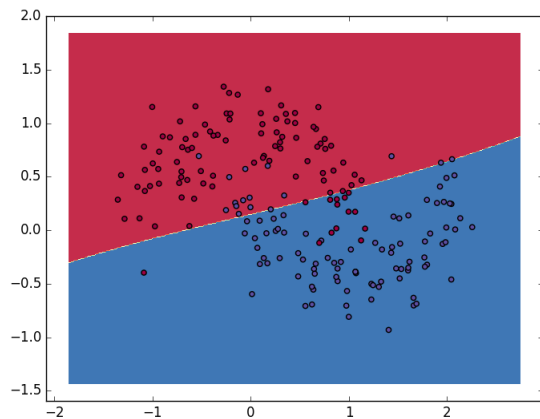


Figure 3: 3 hidden units with sigmoid activation fuctions. Loss = .304

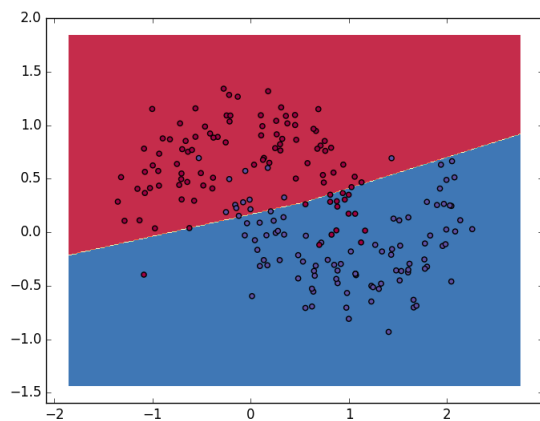


Figure 4: 3 hidden units with ReLU activation fuctions. Loss = .305

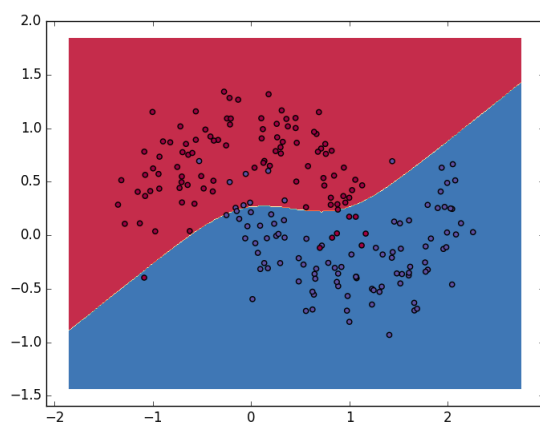


Figure 5: 10 hidden units with tanh activation fuctions. Loss = .246

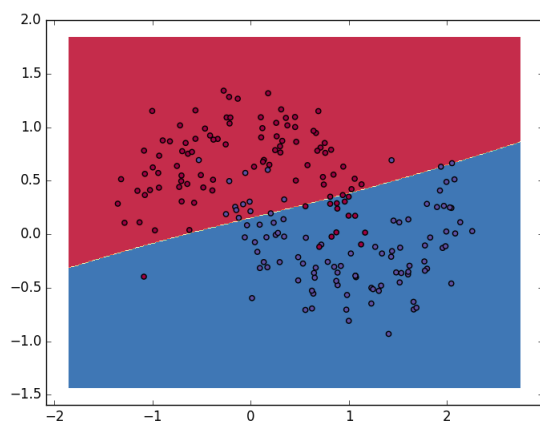


Figure 6: 10 hidden units with sigmoid activation fuctions. Loss = .301

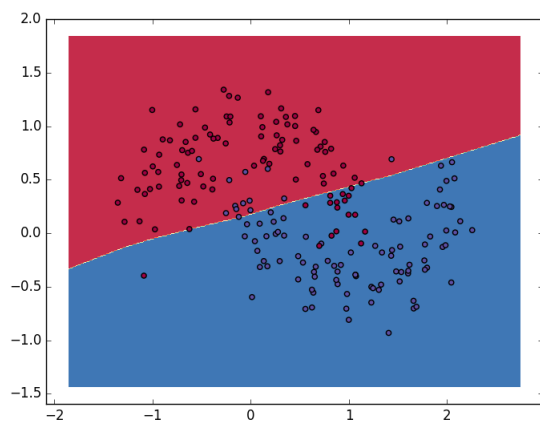


Figure 7: 10 hidden units with ReLU activation fuctions. Loss = .302